DYLAN FERREIRA

DATA ENGINEER | AZURE & MICROSOFT FABRIC | PYSPARK | SQL

California, USA (Open to Remote / Relocation) | 559-205-4802 | ferreira.dylan98@gmail.com | LinkedIn | Portfolio

PROFESSIONAL SUMMARY

Performance-driven Data Engineer with 7+ years of IT experience, including nearly 5 years designing, optimizing, and deploying enterprise data platforms and integration pipelines. Proven record of cutting pipeline runtimes from hours to minutes with Microsoft Fabric, Azure Data Factory (ADF), PySpark, and Delta Lake. Specialized in Medallion architecture (Bronze/Silver/Gold), automated schema reconciliation, and high-throughput ELT/ETL serving distributed multi-campus operations.

TECHNICAL SKILLS

- **Cloud & Data Platforms:** Microsoft Fabric (Lakehouse/Warehouses), Azure Data Factory (ADF), ADLS Gen2, Delta Lake, Azure SQL, SQL Server, PostgreSQL
- **Languages & Scripting:** Python, PySpark, SQL / T-SQL, PowerShell, Bash
- **Data Engineering:** Medallion Architecture (Bronze/Silver), ETL/ELT Pipelines, Incremental Loads, Change Data Capture (CDC), Delta Table Merges, Kimball / Star Schema Modeling, Schema Drift Handling
- **Integration & Tools:** REST APIs, JSON/Parquet, SQLAlchemy, pyodbc, Azure Key Vault, Power BI, Git, Pecan AI

PROFESSIONAL EXPERIENCE

Database Administrator & Data Engineer

San Joaquin Valley College | Jan 2022 – Present | Remote

- Accelerated processing by 95% by migrating 30+ CRM vendor ingestion routines from legacy formats to optimized PySpark in Microsoft Fabric, cutting runtimes from 2+ hours to 5 minutes.
- Architected end-to-end Medallion (Bronze/Silver) pipelines in Azure Data Factory and Microsoft Fabric, processing hundreds of thousands of records per run across 26 campuses.
- Engineered resilient Lakehouse ingestion models that eliminated Power BI reporting disruptions by managing Delta Lake schema drift, concurrent merge conflicts, and datatype inconsistencies.
- Decoupled analytics from core OLTP systems by migrating transactional SQL Server data into ADLS Gen2 and the Microsoft Fabric Lakehouse.
- Automated incremental loads and historical reconciliations, standardizing overlapping month-to-date (MTD) and year-to-date (YTD) datasets for single-source-of-truth accuracy.
- Tuned execution plans, indexes, joins, and partition schemes on heavy analytic and ETL queries, reducing run times from 60+ seconds to sub-second.
- Deployed 4 predictive ML models with Pecan AI for student retention and financial risk, delivering automated scores directly to intervention teams.

Help Desk Specialist, Tier 1

San Joaquin Valley College | Sep 2019 – Dec 2021 | Visalia, CA (Hybrid)

- Managed provisioning, role-based access control, and account governance across Active Directory and Azure AD for 26 campuses.
- Automated administrative provisioning workflows with PowerShell and authored technical documentation adopted department-wide.

IT Field Technician (MTI) & Systems Repair Technician

Jan 2018 – Sep 2019

- Configured, deployed, and supported hardware, network connectivity, and operating systems across commercial client sites.

KEY PROJECTS

- **Canvas LMS Data 2 Pipeline (Fabric / Python / REST):** Built a daily automated pipeline synchronizing 49 LMS relational tables via Python and REST APIs, with reconciliation scripts that validate row counts against upstream APIs.
- **Automated Enterprise Data Ingestion (Azure Key Vault / Document Intelligence):** Built secure micro-pipelines in Python, T-SQL, and Azure Document Intelligence that parse unstructured documents into tabular datasets, with credentials centralized in Azure Key Vault.
- **Automated Pipeline Data Reconciler (Fabric / Power BI):** Built a regression validation tool comparing Fabric pipeline outputs against legacy manual flat-file imports, catching discrepancies before production cutover.

EDUCATION & CERTIFICATIONS

- Associate of Science in Information Technology — San Joaquin Valley College (2017 – 2018)
- Microsoft Certified: Azure Data Fundamentals (DP-900) | Microsoft Certified: Azure Fundamentals (AZ-900)